

# Inference algorithms

# Statistical inference

- Judge the **accuracy** of an estimation or prediction algorithm
  - Efron & Hastie 2016
- **Reliability**
- **Uncertainty**

ISO definition of accuracy: the closeness of a measurement to the true value  
Two components: bias, variance

# Different inference problems

## Estimation

Infer a property of a population (e.g. mean) from a sample

## Model selection

Infer the data generating process from among a set of candidate data-generating processes

## Hypothesis test (association)

Infer that  $y$  is associated with  $x$

## Causation

Infer that  $x$  causes  $y$

Infer the size of an effect due to an experimental intervention (estimation)


Infer that an experimental intervention had an effect (H-test)

## Prediction

Predict the value of a new observation or population state (extrapolation or interpolation)

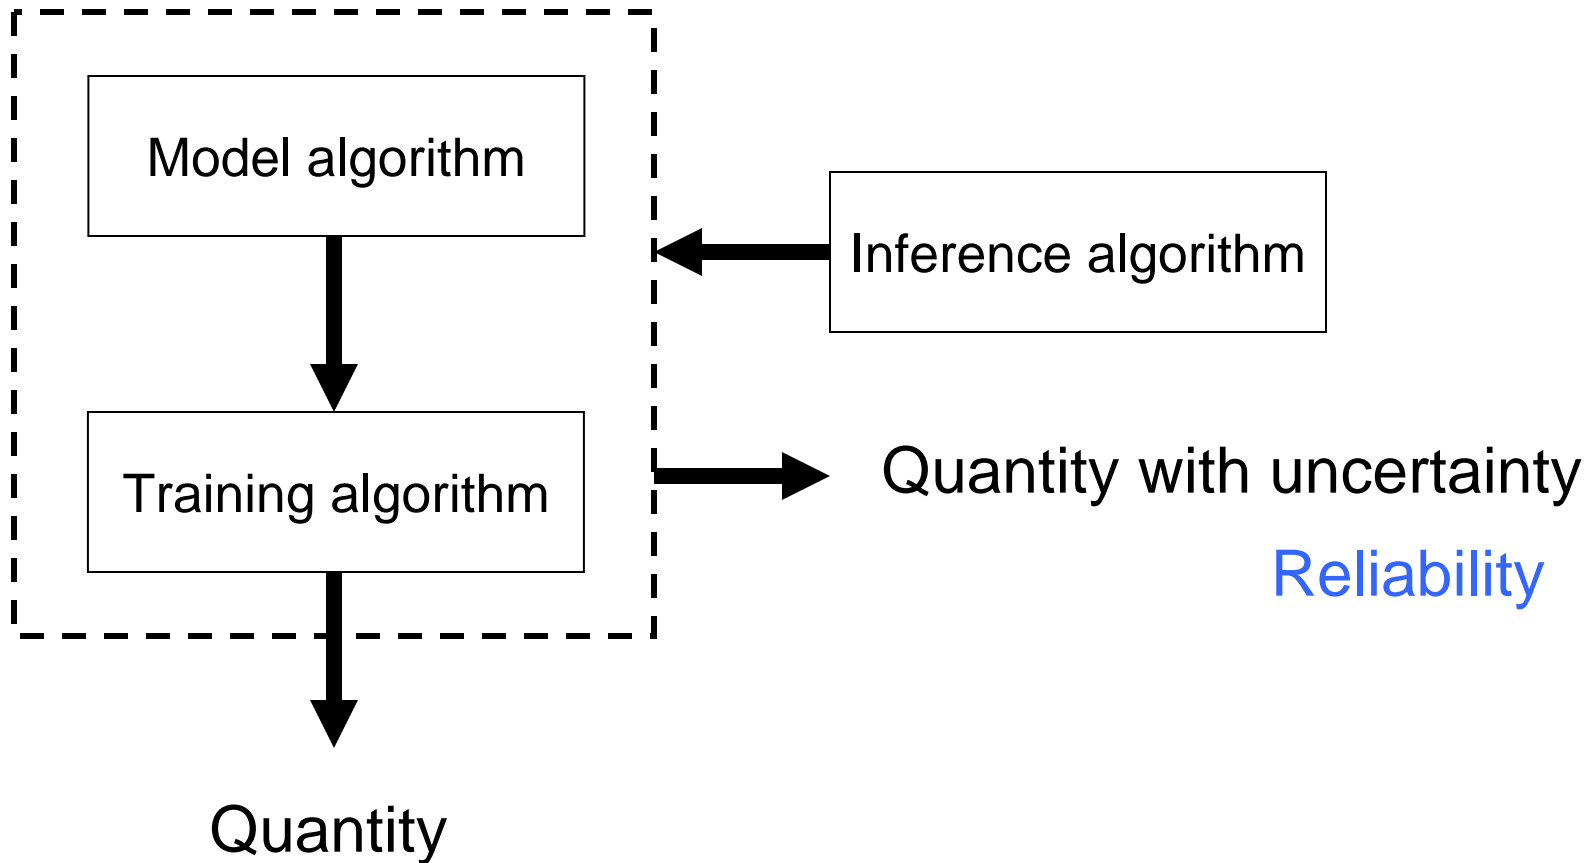
Predict the population state in the future (forecast/extrapolation)

# Algorithms in data science



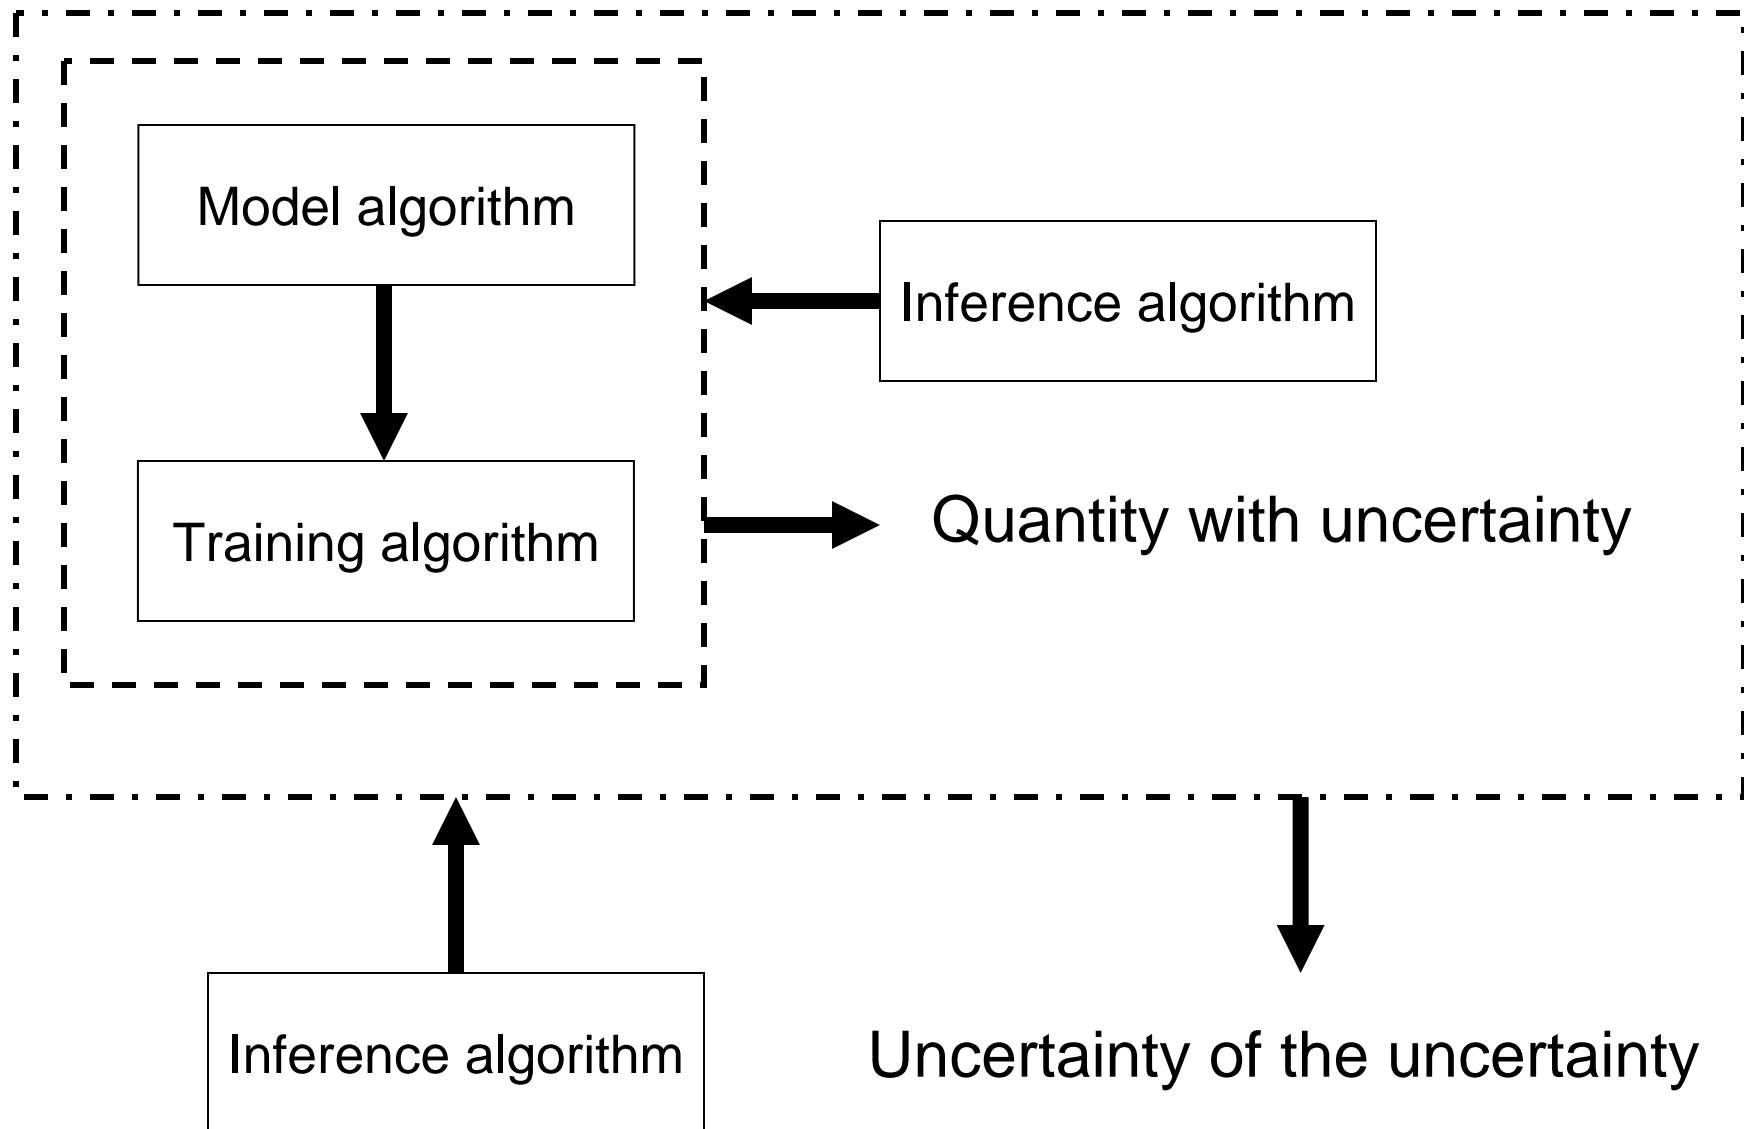
"Dumb" - doesn't say about reliability

# Algorithms in data science



"Dumb" - doesn't say about reliability

# Algorithms in data science



# Inference algorithms

- **Looking back:** considering all the ways data could have happened
- **Looking forward:** predicting new data and testing against them

These are two **big ideas in data science**

# Inference algorithms

- Looking back: considering all the ways data could have happened
  - Frequentist: sampling distribution
  - Bayesian/likelihood:  $P(\text{data}|\text{model})$



# Sampling distribution

132 orange-spotted warblers. 1 indicates female

```
f <- c(1,1,1,1,1,0,0,0,0,0,0,0,0,0,1,1,1,0,1,1,0,1,1,0,0,0,1,1,0,0,1,1,0,1,0,0,0,0,  
       1,1,1,0,1,1,0,1,1,0,0,1,1,0,0,1,1,0,0,0,0,0,0,0,0,0,1,0,0,1,0,0,0,1,0,0,1,1,  
       0,1,0,0,0,1,0,0,0,1,0,0,1,0,0,1,1,0,1,1,0,1,1,0,0,0,0,0,0,0,0,1,0,1,1,1,0,  
       1,0,1,0,0,0,0,0,0,0,1,1,0,0,0,1,1,1,1,0,0,1)
```


Take a sample:

```
sample(f,10)
```

0 1 0 0 0 0 0 1 0 0

sex ratio = 0.2

Our scientific observation



True sex ratio is 0.424

# Sampling distribution algorithm 1

for each possible combination of  $n$  sample units

    sample  $n$  units from the population

    calculate the sample statistic

plot sampling distribution (histogram) of the sample statistic

for bird sex ratio

There are  $3e14$  possible samples.

Too hard! It would take 100 years to compute!

# Sampling distribution algorithm 2

We can invoke the law of large numbers

repeat very many times

- sample  $n$  units from the population

- calculate the sample statistic

- plot sampling distribution (histogram) of the sample statistic

## for bird sex ratio


for a large number of repeated samples

- randomly sample 10 birds from the population

- calculate the sex ratio in the sample

- plot sampling distribution (histogram) of sex ratios

# The sampling distribution for sex ratio




# Confidence interval


- An interval calculated by some procedure that would **contain** (or **cover**) the true population value 95% of the time, **if sampling and calculating an interval were repeated a very large number of times**

Confidence = **reliability** of the **procedure**


True sampling distribution



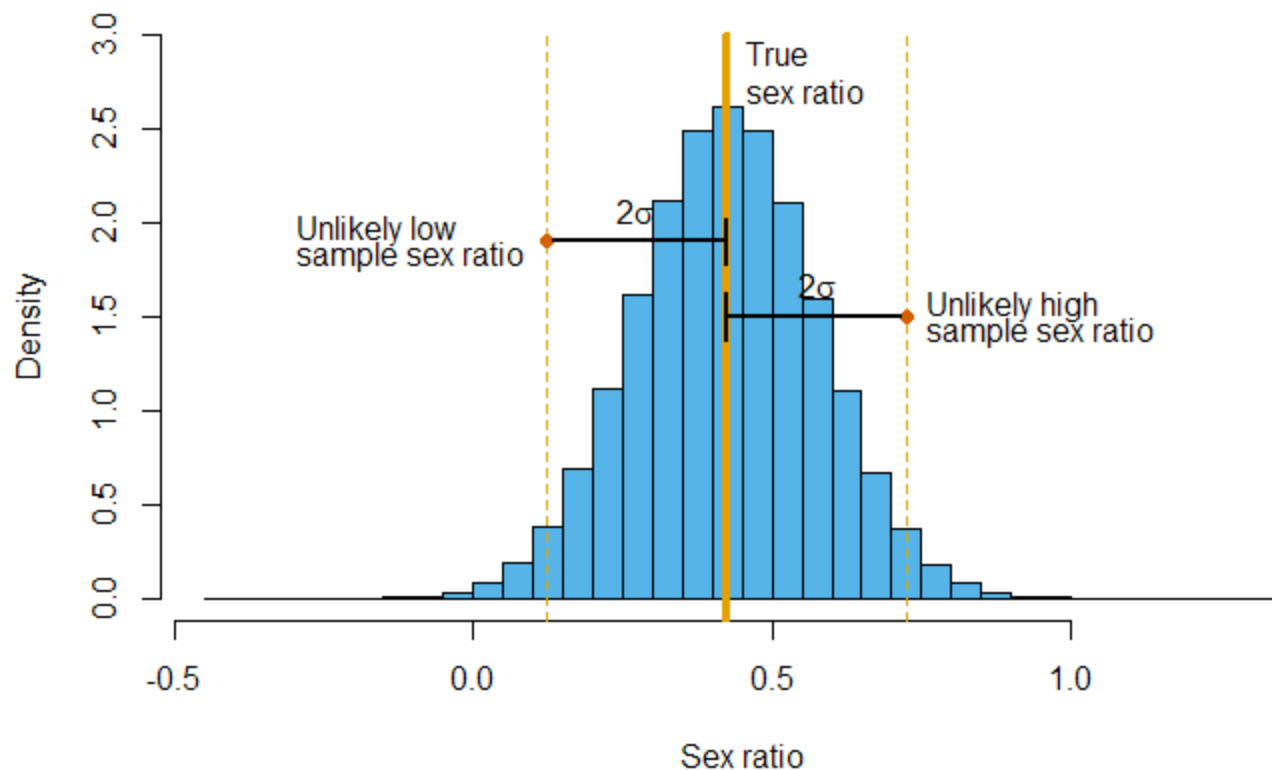
Approximating Normal



Normal overlaid on true



# Using the normal approx for inference





# Plug in principle

- We don't know the true sampling distribution or its parameters
- Plug in the sample instead as an estimate

# Coverage

repeat very many times

- sample  $n$  units from the population


- calculate the sample statistic

- calculate the interval for the sample statistic

- calculate frequency true value is in the interval

**Calibrates** the degree of confidence in the procedure

# First 100 95% confidence intervals



95.6% of the intervals cover the true value  
In first 100, 6 do not cover the true value  
(we expect about 5/100)

# lm() inference algorithms

Sampling distribution for parameters  $\beta_0, \beta_1$

repeat very many times

- sample data from the population

- fit the linear model

- estimate the parameters


plot sampling distribution (histogram) of parameter estimates

Sampling distribution for any other quantities  
(e.g. mean of  $y$  given  $x$ ) is similar


$$y_i = \beta_0 + \beta_1 x_i + e_i$$

Population: normal distribution of errors

Sampling distribution beta\_0



Sampling distribution beta\_1



# Plug-in principle


- We don't have access to the **true** sampling distribution or its parameter values
- **Plug in** the residual standard error from the **sample** to estimate the **parameters** ( $\sigma$ ) of the **sampling distribution**

# P-values

- The probability of a sample statistic as large **or larger** than the one observed **given that some hypothesis is true**
- Obtained from the **sampling distribution** of the parameters ( $t$  standardized)
- $t$  is  $\beta$  in standard error units

# Confidence vs prediction intervals

CI: uncertainty in mean response



PI: uncertainty  
in individual  
response



# Robustness

- Normality of  $e_i$  is not that crucial
- **More relevant:** sampling distributions for  $\beta$  are Normal
  - central limit theorem says whatever the  $e_i$ s, the sampling distribution will tend Normal
- Most problematic: when  $e_i$  is asymmetrical or heteroscedastic

# R code - most common inferences

```
plot(x,y)
fit <- lm(y ~ x)
summary(fit)
confint(fit)
newd <- data.frame(x = seq(min(x), max(x), length.out=100))
pred_w_ci <- cbind(newd,predict(fit, newd, interval = "confidence"))
pred_w_pi <- cbind(newd,predict(fit, newd, interval = "prediction"))
lines(pred_w_ci[,c(1,nrow(pred_w_ci))],c("x","fit"),col="#56B4E9")
lines(pred_w_ci[,c("x","lwr")],col="grey")
lines(pred_w_ci[,c("x","upr")],col="grey")
lines(pred_w_pi[,c("x","lwr")],col="grey",lty=2)
lines(pred_w_pi[,c("x","upr")],col="grey",lty=2)
plot(fit,1:6)
```